

Stop scoring the resume. Score the proof.

Redrob's Resume Ranker grades what a candidate *wrote about themselves*. That's why the data is full of keyword-stuffers and ~80 impossible honeypots.

Pramaan (प्रमाण – "proof") doesn't read the profile. It verifies the **evidence the profile points at**: real commits, live deployments, published work – inside the claimed time window.

Fraud isn't a text problem. It's a provenance problem. Provenance is cheap to verify, expensive to fake.

hack2skill × Redrob "India Runs" Ideathon – Sub-track 1 (deep technical / AI-native)

Where Pramaan fits

Redrob's promise: hire on *real data*, not resumes. **Redrob's reality:** the Resume Ranker still scores resume *text* – so it ranks the best self-describers, not the best builders.

The signal Redrob already exposed:

- Profiles stuffed with keywords to game the ranker
- ~80 "impossible" honeypot profiles seeded to catch fakes

Ideathon scope (Sub-track 1): an AI-native, explainable PoC Redrob could fold into its roadmap.

Pramaan is the missing piece between a claim and a hire: a **verification layer** that sits *under* the Resume Ranker, not beside it. We don't replace the ranker. We feed it ground truth.

The problem: you're grading self-description

A resume says **"React, 4 years."** The ranker scores the *sentence*.

Three ways that sentence lies – all invisible to a text scorer:

Failure	What the text says	What's actually true
Keyword-stuffer	"React, Node, K8s, Rust, ML"	Zero artifacts behind any of it
Honeypot / fabricator	"Led X in 2021"	Timeline doesn't reconcile – couldn't have
Borrowed credit	"Built this repo"	Forked it; never authored a line

The harder you tune a text model, the better it ranks **fluent liars**. You can't out-prompt this. The ground truth isn't in the profile – it's in the **artifacts the profile cites**.

Opportunity & vision

Today: hiring trust is manual. A senior engineer opens the GitHub, reads commits, checks if the deploy is live. It doesn't scale to Redrob's funnel.

The vision: make that senior-engineer judgment **automatic, cited, and explainable** – one Evidence Card per claimed skill.

Praamaan turns every skill claim into a verdict: **verified · unverified · contradicted** – each with the artifacts it's based on and a plain-language reason a recruiter can read in 5 seconds.

The output is a **Trust Layer**, not a ranking. Redrob keeps owning the ranking decision; we give it evidence it can defend to a hiring manager.


Solution: the Proof-of-Work Graph

For every claimed skill, Pramaan runs one agent loop:

1. **RESOLVE** – find the candidate's real artifacts: GitHub handle, deployed URLs, blog
2. **FETCH** – commit metadata + diffs (GitHub API), URL liveness (Playwright), published writing (Jina Reader), discovery (Exa)
3. **RECONCILE** – the LLM reasons over *retrieved evidence only* – never the profile text – to decide if the claim holds in its window
4. **ATTRIBUTE** – forked-vs-authored · AI-generated-code **flag** (never auto-reject) · timeline reconciliation (the honeypot kill-shot)
5. **EMIT** – an **Evidence Card**: `{ skill, verdict, confidence, cited_artifacts[], plain_language_reason }`

Core rule: the model is allowed to see proof. It is *never* allowed to see the candidate's own description of themselves.


User journey



Recruiter pastes a profile; Pramaan returns cited Evidence Cards in seconds.

Candidate submits Redrob profile + linked artifacts (GitHub, deploys, writing). **Pramaan** runs the loop per skill, in the background – no candidate friction. **Recruiter** opens the profile and sees Evidence Cards inline:

AI logic & decision flow




Verdict is deterministic on evidence; the LLM only writes the reason.

The verdict is **deterministic on evidence, reasoned by the LLM**:

- **No artifacts found** → UNVERIFIED (keyword-stuffer). We never punish; we just don't confirm.
- **Artifacts exist, authored, within window** → VERIFIED + citations.
- **Artifacts contradict the claim** (forked-only, or timeline doesn't reconcile) → CONTRADICTED (honeypot caught).

System architecture

Data, context & intelligence layer



Retrieval-grounded: the model reasons over fetched evidence, not profile text.

Praamaan's intelligence is **retrieval-grounded** – the model reasons over evidence we fetched, not over training memory or profile text.

- **Evidence corpus (per candidate, ephemeral):** commit diffs, repo ownership graph, deploy responses, published articles
- **Window filter:** every artifact is timestamped and checked against the *claimed* window – this is what kills the honeypot
- **Attribution graph:** author vs contributor vs forker, resolved from commit authorship not repo ownership
- **Citation store:** every verdict carries the artifacts it stands on, so nothing is unexplainable

No candidate PII is stored. We keep verdicts + citations, then discard raw evidence.

Scalability & feasibility: cheap by design


The expensive step is *fetching*, and fetching is **mostly free**:

- GitHub API, URL liveness, Jina Reader – public, near-zero cost
- The LLM only sees a **distilled evidence summary**, not raw diffs → small context → cheap inference
- **OpenRouter** lets us route easy verdicts to a small/cheap model and escalate only ambiguous ones
- Cache by `{handle, skill, artifact-hash}` – a candidate re-checked tomorrow costs ₹0

Target: **a few cents per fully-cited candidate**. Verification stays cheaper than fraud – that's the whole point.

Serverless on Vercel means it scales to Redrob's funnel with zero ops.

Redrob ecosystem integration



A trust layer beneath the Resume Ranker; the Verified badge compounds the network effect.

Pramaan slots in as a **trust layer beneath the Resume Ranker** – a clean contract, not a rebuild:

- Redrob sends `{ candidate_id, claimed_skills[], artifact_links[] }`
- Pramaan returns `Evidence Cards[]`
- The Ranker uses verdicts as **features** (verified skills weighted up, contradicted ones flagged for review)

Impact & what we measure

We frame these as **targets we instrument**, not claims:

- **Honeypot catch rate** – % of the ~80 seeded fakes marked **CONTRADICTED** . North-star metric.
- **Keyword-stuffer precision** – claims with zero artifacts correctly marked **UNVERIFIED** .
- **Recruiter time-to-trust** – minutes of manual GitHub digging removed per profile.
- **False-flag rate** – legit candidates wrongly contradicted (must stay near zero; we bias toward **UNVERIFIED** over false **CONTRADICTED**).
- **Cost per verified candidate** – kept in cents.

Every number is auditable because every verdict cites its artifacts.

Bharat: proof without a resume

Most of India's workforce has **no GitHub and no resume** – but they have proof of work.

Pramaan on vernacular voice notes: a candidate describes a job in Tamil/Hindi/Telugu; we transcribe, then verify against *real artifacts* – a UPI transaction history, a delivery-app rating, photos of completed work, a shop's Google listing.

Same engine, different evidence connectors. The thesis generalizes: **score the proof, whatever form it takes.**

Verified Skill Passport: a portable, cited proof-of-work record a worker carries across platforms. Owned by the worker, verifiable by anyone. That's the long-term moat – and it's India-first.

Roadmap

Now (PoC): GitHub + deploy + writing → 3 Evidence Card verdicts, live.


Next: ranker integration contract · caching layer · confidence calibration against the honeypot set.

Then: more evidence connectors (npm, Kaggle, design portfolios, Stack Overflow) · recruiter feedback loop to tune confidence.

Bharat track: vernacular voice intake · non-code evidence connectors · Verified Skill Passport pilot.

Each step adds **depth on the core verdict**, not surface area.

Demo



Recruiter pastes a profile; Pramaan returns cited Evidence Cards in seconds.

Live, 60 seconds: paste a GitHub handle + claimed skills → three Evidence Cards render side by side.

- **● VERIFIED** – "React" → cites 3 authored commits + a live deployed URL, all inside the claimed window

The asymmetry is the moat

Faking a resume is free. Faking a year of authored commits, a live product, and a reconcilable timeline is not.

Pramaan stops scoring **what candidates say** and starts scoring **what they can prove** – cited, explainable, cheap, and India-ready.

Stop grading self-description. Verify the proof.

प्रमाण – *Pramaan*

Thank you. Questions?